

# COMS4040A & COMS7045A: High Performance Computing & Scientific Data Management Introduction to OpenMP: Part II

Hairong Wang

School of Computer Science,  
University of the Witwatersrand, Johannesburg

2019-2-28

- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - Single Worksharing Construct
  - Task Worksharing Construct
- 2 Data Environment
  - Changing Storage Attributes

- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - Single Worksharing Construct
  - Task Worksharing Construct

- 2 Data Environment
  - Changing Storage Attributes

- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - Single Worksharing Construct
  - Task Worksharing Construct
- 2 Data Environment
  - Changing Storage Attributes

# A Review on OpenMP

The OpenMP constructs we've learned so far:

- To create a team of threads (or a parallel region):
- To share work between threads:
- To prevent race conditions (or conflicts):
- Data environment (or data scoping) clauses:

# A Review on OpenMP Contd.

```
1 double compute_pi(double step){
2     double pi, x, sum=0.0;
3     omp_set_num_threads(NUM_THREADS);
4     #pragma omp parallel
5     {
6         #pragma omp for reduction(+:sum) private(x)
7         for (int i=1;i<= num_steps; i++){
8             x = (i-0.5)*step;
9             sum = sum + 4.0/(1.0+x*x);
10        }
11    }
12    pi = step * sum;
13    return pi;
14 }
```

# A Review on OpenMP Contd.

```
1 #pragma omp for [clause[,] clause]...  
2 {...}
```

- Possible clauses

- private, firstprivate, lastprivate
- reduction
- schedule, nowait, collapsed, ordered

- Implicit barrier at the end of `for` loop

- `nowait`

- modifies a `for` directive
- avoids implicit barrier at end of `for`

- `ordered`

- `for` loops with carried dependence
- The enclosed block of code is executed in the order in which iterations would be executed sequentially.

## The collapse Clause

- The iterations of the  $k$  and  $j$  loops are collapsed into one loop. and that loop is then divided among the threads in the current team.

```
1 void work(int a, int j, int k);
2 void main()
3 {
4     int j, k, a;
5     #pragma omp parallel num_threads(2)
6     {
7         #pragma omp for collapse(2) ordered private(j,k) schedule
            (static,3)
8         for (k=0; k<3; k++)
9             for (j=0; j<2; j++)
10                {
11                    #pragma omp ordered
12                    printf("%d %d %d\n", omp_get_thread_num(), k, j);
13                    /* end ordered */
14                    work(a,j,k);
15                }
16     }
17 }
```



## The collapse Clause

- The iterations of the  $k$  and  $j$  loops are collapsed into one loop. and that loop is then divided among the threads in the current team.

```
1 void work(int a, int j, int k);
2 void main()
3 {
4     int i, j, k, a;
5     #pragma omp parallel num_threads(2)
6     {
7         #pragma omp for collapse(2) private(i,j,k)
8         for (k=0; k<3; k++)
9             for (j=0; j<2; j++)
10                for (i = 0; i<4; i++)
11                    {
12                        printf("%d %d %d\n", i, k, j);
13                        work(a,j,k);
14                    }
15
16     }
17 }
```

# More on for construct

- OpenMP parallelizes `for` loops that are in canonical form.
- Loops in canonical form take one of the following forms.

```
                                index++
                                ++index
                                index < end   index--
                                index <= end  --index
for(index=start; index >= end; index += incr  )
                                index > end   index -= incr
                                index=index+incr
                                index=incr+index
                                index=index-incr
```

- 1 Worksharing
  - A Review on OpenMP
  - **Sections/section construct**
  - Single Worksharing Construct
  - Task Worksharing Construct
- 2 Data Environment
  - Changing Storage Attributes

# Sections/Section Construct

- `sections` directive enables specification of **task parallelism**
- The `sections` worksharing construct gives a different structured block to each thread.
- Syntax:

```
1  #pragma omp sections [clause[[,] clause]...]
2  {
3      [#pragma omp section]
4          structured block
5      [#pragma omp section]
6          structured block]
7      ...
8  }
```

- **clauses:** `private`, `firstprivate`, `lastprivate`, `reduction`, `nowait`
- Each section must be a structured block of code that is independent of other sections.
- There is an implicit barrier at the end of a sections construct

# Examples — firstprivate

```
1 #include <omp.h>
2 #include <stdio.h>
3 #define NT 4
4 int main( ) {
5     int section_count = 0;
6     omp_set_dynamic(0);
7     omp_set_num_threads(NT);
8     #pragma omp parallel
9     #pragma omp sections firstprivate( section_count )
10 {
11     #pragma omp section
12     {
13         section_count++;
14         printf( "section_count %d\n", section_count );
15     }
16     #pragma omp section
17     {
18         section_count++;
19         printf( "section_count %d\n", section_count );
20     }
21 }
22 return 0;
23 }
```

# Examples — lastprivate clause

```
1 void lastpriv (int n, float *a, float *b)
2 {
3     int i;
4     .....
5     #pragma omp parallel
6     {
7         #pragma omp for lastprivate(i)
8         for (i=0; i<n-1; i++)
9             a[i] = b[i] + b[i+1];
10    }
11
12    a[i]=b[i];
13 }
```

# Examples — nowait clause

```
1 #include <math.h>
2 void nowait_example2(int n, float *a, float *b, float *c, float
   *y, float *z)
3 {
4     int i;
5     #pragma omp parallel
6     {
7         #pragma omp for schedule(static) nowait
8         for (i=0; i<n; i++)
9             c[i] = (a[i] + b[i]) / 2.0f;
10        #pragma omp for schedule(static) nowait
11        for (i=0; i<n; i++)
12            z[i] = sqrtf(c[i]);
13        #pragma omp for schedule(static) nowait
14        for (i=1; i<=n; i++)
15            y[i] = z[i-1] + a[i];
16    }
17 }
```



## Example (1)

Parallelize the sequential *quicksort* program (`qsort_v00.c`) using OpenMP `sections` construct.



# Quick Sort Algorithm

```
1 q_sort(left, right, data){
2     q = partition(left, right, data);
3     q_sort(left, q-1, data);
4     q_sort(q+1, right, data);
5 }
6 partition(left, right, data){
7     x = data[left];
8     p = left+1; r = right;
9     while (p < r){
10         while (data[r] > x)
11             r=r-1;
12         while (data[p] <= x && p < r)
13             p=p+1;
14         if (p <= r)
15             swap(data[p], data[r]);
16             p=p+1; r=r-1;
17     }
18     swap(data[left], data[r]);
19     return r;
20 }
```

- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - **Single Worksharing Construct**
  - Task Worksharing Construct
- 2 Data Environment
  - Changing Storage Attributes

# Single Worksharing Construct

- The *single* construct denotes a block of code that is executed by only one thread.
- Syntax:

```
1  #pragma omp single [clause[[,] clause]...]  
2      structured block
```

- **clauses:** `private`, `firstprivate`, `copyprivate`, `nowait`
- A barrier is implied at the end of the *single* block, unless a `nowait` clause is specified.
- This construct is ideally suited for I/O or initialization.

# Single Worksharing Construct Contd.

```
1 void work1() {}
2 void work2() {}
3 void single_example()
4 {
5     #pragma omp parallel
6     {
7         #pragma omp single
8         printf("Beginning work1.\n");
9         work1();
10        #pragma omp single
11        printf("Finishing work1.\n");
12        #pragma omp single nowait
13        printf("Finished work1 and beginning work2.\n");
14        work2();
15    }
16 }
```

**Exercise 1:** For simple worksharing examples, compile and run “worksharing1.c”.

# Master construct

`master construct`:

- The `master` construct specifies a structured block that is executed by the master thread of the team.
- There is no implied barrier either on entry to, or exit from, the `master` construct.

# Combined Parallel Work-Sharing Constructs

Combined parallel worksharing constructs are shortcuts that can be used when a parallel region comprises precisely one worksharing construct.

```
1 //Full for and sections
  versions
2 #pragma omp parallel
3 {
4     #pragma omp for
5     for-loop
6 }
7 #pragma omp parallel
8 {
9     #pragma omp sections
10    {
11        [#pragma omp section]
12        structured block
13        [#pragma omp section
14        structured block]
15        ...
16    }
17 }
```

```
1 //Combined for version
2 #pragma omp parallel for
3 {
4     for-loop
5 }
6 //Combined sections version
7 #pragma omp parallel sections
8 {
9     [#pragma omp section]
10    structured block
11    [#pragma omp section
12    structured block]
13    ...
14 }
15 }
```



- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - Single Worksharing Construct
  - Task Worksharing Construct

- 2 Data Environment
  - Changing Storage Attributes

# Task Worksharing Construct

- Tasks are independent units of work.
- Threads are assigned to perform the work of each task.



## Task Construct Syntax

```
#pragma omp task [clause[,] clause]...  
structured block
```

where clause can be

- `if(expression)`: if *expression*=TRUE, then the task is immediately executed.
- `shared`
- `private`
- `firstprivate`
- `default( shared|none )`
- `untied`

# Task Construct Contd.

- Two activities: *packaging* and *execution*
  - Each encountering thread packages a new instance of task
  - Some thread in the team executes the task at some time later or immediately.

# Task Construct Contd.

- Two activities: *packaging* and *execution*
  - Each encountering thread packages a new instance of task
  - Some thread in the team executes the task at some time later or immediately.
- Task barrier: The **taskwait** directive:

# Task Construct Example

## Example (2)

```
1 .....  
2 #pragma omp parallel  
3 {  
4     #pragma omp single private (p)  
5     {  
6         p=list_head;  
7         while (p) {  
8             #pragma omp task  
9                 processwork (p);  
10            p=p->next;  
12        }  
13    }  
14 }
```

# Task Construct Contd.

## When tasks are guaranteed to be completed?

- At thread or task barriers
- At the directive: `#pragma omp barrier`
- At the directive: `#pragma omp taskwait`

## Example (3)

```
1 #pragma omp parallel
2 {
3     #pragma omp task
4     foo();
5     #pragma omp barrier
6     #pragma omp single
7     {
8         #pragma omp task
9         bar();
10    }
11 }
```

# Task Construct Example

## Example (4)

### Understanding Task Construct

```
1 int main(int argc, char *argv[]){  
3     printf("A ");  
4     printf("race ");  
5     printf("car ");  
  
7     printf("\n");  
8     return 0;  
9 }
```

# Task Construct Example

## Example (5)

### Understanding Task Construct

```
1 int main(int argc, char *argv[]){  
2     #pragma omp parallel num_threads(2)  
3     {  
4         printf("A ");  
5         printf("race ");  
6         printf("car ");  
7     }  
8     printf("\n");  
9     return 0;  
10 }
```

# Task Construct Example

## Example (6)

### Understanding Task Construct

```
1 int main(int argc, char *argv[]){  
2     #pragma omp parallel  
3     {  
4         #pragma omp single  
5         {  
6             printf("A ");  
7             printf("race ");  
8             printf("car ");  
9         }  
10    }  
11    printf("\n");  
12    return 0;  
13 }
```



# Task Construct Example

## Example (8)

### Understanding Task Construct

```
1 int main(int argc, char *argv[]){  
2     #pragma omp parallel  
3     {  
4         #pragma omp single  
5         {  
6             printf("A ");  
7             #pragma omp task  
8             printf("race ");  
9             #pragma omp task  
10            printf("car ");  
11        }  
12    }  
13    printf("\n");  
14    return 0;  
15 }
```

# Task Construct Example

## Example (9)

### Understanding Task Construct

```
1 int main(int argc, char *argv[]){  
2     #pragma omp parallel  
3     {  
4         #pragma omp single  
5         {  
6             printf("A ");  
7             #pragma omp task  
8             printf("race ");  
9             #pragma omp task  
10            printf("car ");  
11            printf("is fun to watch ");  
12        }  
13    }  
14    printf("\n");  
15    return 0;  
16 }
```

# Task Construct Example

## Example (10)

### Understanding Task Construct

```
1 int main(int argc, char *argv[]){  
2     #pragma omp parallel  
3     {  
4         #pragma omp single  
5         {  
6             printf("A ");  
7             #pragma omp task  
8             printf("race ");  
9             #pragma omp task  
10            printf("car ");  
11            #pragma omp taskwait  
12            printf("is fun to watch ");  
13        }  
14    }  
15    printf("\n");  
16    return 0;  
17 }
```

# Task Construct Example

## Example (11)

### Tree traversal using task

```
2 void traverse(node *p) {  
3     if (p->left)  
4         #pragma omp task  
5         traverse(p->left);  
6     if (p->right)  
7         #pragma omp task  
8         traverse(p->right);  
9     process(p->data);  
10 }
```

# Task Construct Example

## Example (12)

### Tree traversal using task

```
2 void traverse(node *p) {  
3     if (p->left)  
4         #pragma omp task  
5         traverse(p->left)  
6     if (p->right)  
7         #pragma omp task  
8         traverse(p->right)  
9     #pragma omp taskwait  
10    process(p->data);  
11 }
```

## Exercise 1

Parallelize the program `fib_v00.c` for computing the  $n$ th Fibonacci number.

# Task Construct Contd.

## Task switching: *untied*:

```
1 #define ONEBILLION 10000000000L
2 .....
3 #pragma omp parallel
4 {
5     #pragma omp single
6     {
7         for(i=0; i<ONEBILLION; i++)
8             #pragma omp task
9                 process(item[i]);
10    }
11    .....
12    /* Untied task: any other thread is eligible to resume
13       the task generating loop*/
14    #pragma omp single
15    {
16        #pragma omp task untied
17        for(i=0; i<ONEBILLION; i++)
18            #pragma omp task
19                process(item[i]);
20    }
21 }
```

## Exercise 2:

Parallelize the sequential *quicksort* program (`qsort_v00.c`) using OpenMP `task` construct. Compare the performances of three versions of quicksort implementations, i.e., the sequential implementation, the OpenMP implementation using `sections` construct, and the OpenMP implementation using `task` construct.

# Summary for clause applicability

Clause	Directive					
	PARALLEL	DO/for	SECTIONS	SINGLE	PARALLEL DO/for	PARALLEL SECTIONS
IF	●				●	●
PRIVATE	●	●	●	●	●	●
SHARED	●	●			●	●
DEFAULT	●				●	●
FIRSTPRIVATE	●	●	●	●	●	●
LASTPRIVATE		●	●		●	●
REDUCTION	●	●	●		●	●
COPYIN	●				●	●
SCHEDULE		●			●	
ORDERED		●			●	
NOWAIT		●	●	●		

Figure: Clause applicability for OpenMP constructs



- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - Single Worksharing Construct
  - Task Worksharing Construct

- 2 Data Environment
  - Changing Storage Attributes

- Shared memory programming model. Hence most variables are shared by default.
- Global variables are shared among threads. For C, file scope variables and static variables.
- Not everything is shared. Loop indexes are `private` Automatic variables within a statement block are `private`.

## Example

What are the data scoping characteristics of the variables in the following code fragment?

```
1 .....  
2 int b[3], i;  
3 char *ptr;  
4 cptr=malloc(sizeof(char)*32);  
5 #pragma omp parallel  
6 {  
7     #pragma omp for  
8     for (i=0; i<3; i++)  
9         b[i]=i;  
10 }
```

- 1 Worksharing
  - A Review on OpenMP
  - Sections/section construct
  - Single Worksharing Construct
  - Task Worksharing Construct

- 2 Data Environment
  - Changing Storage Attributes

# Changing Storage Attributes

- Changing storage attributes using clauses:
  - shared
  - private: private(var)

## Example

What is the result of the following code being executed?

```
1 .....  
2 int tmp=0;  
3 #pragma omp parallel for  
   private (tmp)  
4   for(int j=0; j<1000; ++j)  
5     tmp+=j;  
6 printf("%d\n", tmp);
```

- firstprivate: Variables are initialized from shared variable.

## Example (6)

What is the result of the following code being executed?

```
1 .....  
2 int tmp=0;  
3 /*Each thread gets its own  
   copy of tmp with an  
   initial value of 0*/  
4 #pragma omp parallel for  
   firstprivate(tmp)  
5   for(int j=0; j<1000; ++j){  
6     if ((j%2)==0) tmp++;  
7     A[j]=tmp;  
8   }
```

## Changing Storage Attributes Contd.

- *default* clause: The *default* clause allows the user to specify a default scope for all variables in a parallel region.
- For C/C++, the syntax is *default(shared | none)*.

### Example

Consider the following variables:

```
int A=1, B=1, C=1;  
#pragma omp parallel private(B) firstprivate(C)
```

Questions:

- Are *A*, *B*, *C* local to each thread or shared inside the parallel region?
- What are their initial values inside and values after the parallel region?

# Changing Storage Attributes Contd.

- lastprivate: Variables update shared variable using value from last iteration.

## Example

```
1 void sq2(int n, double *lastterm)
2 {
3     double x; int i;
4     #pragma omp parallel for lastprivate(x)
5     for(int i=0; i<1000; i++){
6         x=a[i]*a[i]+b[i]*b[i];
7         b[i]=sqrt(x);
8     }
9     /*x has the value it held for the last sequential iteration,
10      i.e., for i=(n-1)*/
11     *lastterm = x;
12 }
```

# More Synchronization Constructs

- **#pragma omp master**

The *master* construct denotes a structured block that is only executed by the master thread. Other threads just skip it.

```
1 #pragma omp parallel shared(a,b) private(i)
2 {
3     #pragma omp master
4     {
5         a = 10;
6         printf("Master construct is executed by thread %d\n",
7             omp_get_thread_num());
8     }
9     #pragma omp barrier
10    #pragma omp for
11    for (i=0; i<n; i++)
12        b[i] = a;
13 } /*-- End of parallel region --*/
14
15 printf("After the parallel region:\n");
16 for (i=0; i<n; i++)
17     printf("b[%d] = %d\n",i,b[i]);
18 }
```



# More Synchronization Constructs

- **nowait** clause: Removes the implicit barrier at the end of a worksharing construct.

## Example

```
1 #pragma omp parallel shared(A,B,C) private(id)
2 {
3     id=omp_get_thread_num();
4     A[id]=big_calc1(id);
5     #pragma omp barrier
6     #pragma omp for
7         for (i=0; i<N; i++) {
8             C[i]=big_calc2(i,A);
9         } //Implicit barrier
10    #pragma omp for nowait
11        for (i=0; i< N; i++) {
12            B[i]=big_calc3(C,i);
13        } //No implicit barrier due to nowait
14    A[id]=big_calc4(id);
15 } //Implicit barrier at the end of a parallel region
```

# More OpenMP Clauses: Copyprivate Clause

- Used with a *single* region (only) to broadcast values of privates from one member of a team to the rest of the team.

## Example

Using *copyprivate* clause.

```
1 #include <omp.h>
2 void input_parameters(int, int);
3 void do_work(int, int);
4 void main() {
5     int Nsize, choice;
6     #pragma omp parallel private(Nsize, choice)
7     {
8         #pragma omp single copyprivate(Nsize, choice)
9         input_parameters(Nsize, choice);
10        do_work(Nsize, choice);
11    }
12    .....
13 }
```

# Loop Carried Dependency

- Loop carried dependency: Dependencies between instructions in different iterations of a loop;
- What are the dependencies in the following loop?

```
for (i=0; i<N; i++)  
{  
  B[i]=tmp;  
  A[i+1]=B[i+1];  
  tmp=A[i];  
}
```

# Loop Carried Dependency Contd.

It helps to unroll the loop to see the dependencies.

```
1  i=1:
2      B[1]=tmp;
3      A[2]=B[2];
4      tmp=A[1];

6  i=2:
7      B[2]=tmp;
8      A[3]=B[3];
9      tmp=A[2];

11 i=3:
12     B[3]=tmp;
13     A[4]=B[4];
14     tmp=A[3];
15     . . . . .
```

# Loop Carried Dependency Example

```
1 //Loop dependency
2 int i, j, A[MAX];
3 j=5;
4 for (i=0; i<MAX; i++) {
5     j+=2;
6     A[i]=big(j);
7 }
8 .....
```

# Race Condition

- Shared data requires special care.
- A problem may arise in case multiple threads access the same memory section simultaneously.
- This can lead to a **race condition**.
  - Read only data is no problem
  - Updates have to be checked for race condition.
- Different runs of same program might give different results due to race conditions.

# Lab Exercises

Complete all the exercises in the class.

# References

- Introduction to Parallel Computing. Ananth Grama, Anshul Gupta, George Karypis, Vipin Kumar.  
<https://www.cs.purdue.edu/homes/ayg/book/Slides/>
- Using OpenMP: Portable Shared Memory Parallel Programming (Scientific and Engineering Computation), by Barbara Chapman, Gabriele Jost and Ruud van der Pas. The MIT Press, 2007.
- <https://computing.llnl.gov/tutorials/openMP/#Introduction>
- OpenMP Application Programming Interface,  
<https://www.openmp.org/resources/refguides/>
- OpenMP Application Programming Interface Examples,  
<https://www.openmp.org/specifications/>