Machine Learning – COMS3007

Introduction to Machine Learning

Benjamin Rosman
Benjamin.Rosman1@wits.ac.za / benjros@gmail.com

February 2018

Course Details

- Lecturers:
  - Dr Benjamin Rosman / Mr Phumlani Khoza
- Contact details:
  - Offices: MSB UG09 / MSB 145
  - Email: Benjamin.Rosman1@wits.ac.za / Phumlani.Khoza@wits.ac.za
- Lecture Venues and Times:
  - All lectures are in G201 (Geology Building) on Wednesdays from 08h00 - 09h45.
- Tutorials and Labs:
  - There is a tutorial/laboratory session on Mondays from 14h15 - 17h00 in MSL004.
- Consultation Times:
  - Mondays: 13h15 - 14h15

Assessments

- Labs/assignments: 20%
  - Weekly labs, due at the beginning of the next week's lab session
  - One assignment
- Test: 20%
  - There will be a test in a lab session towards the end of the first block (date, venue and scope to be announced)
- Exam: 60%

Course Outline

1. Introduction to Machine Learning
2. Optimisation
3. Naïve Bayes and Probability
4. Decision Trees
5. Linear Regression
6. Logistic Regression
7. Neural Networks
8. k-Nearest Neighbours
9. K-means Clustering
10. Principal Components Analysis
11. Reinforcement Learning
12. Practical Application of ML Methods

Textbooks

There is no prescribed textbook. We will loosely be following these textbooks:

- Machine Learning, Tom M. Mitchell, McGraw-Hill, 1997
- Reinforcement Learning: An Introduction, Richard S. Sutton and Andrew G. Barto, The MIT press, 1998.
- Numerical Analysis, Richard L. Burden and J. Douglas Faires, Brooks/Cole, Cengage Learning, 2011.
- Numerical Optimization, Jorge Nocedal and Stephen J. Wright, Springer-Verlag, 1999.

There are many texts on Machine Learning in the library and on the web that have information on the above topics.

You can also look up the Coursera course “Machine Learning” by Andrew Ng (Stanford University) which follows very similar content.

Programming Languages

All labs and assignments will require programming.

We recommend using Python, but you may use Matlab if you prefer.

Any issues caused by your choice of programming language are largely yours to deal with.

Mathematics

This is a very maths-heavy course!

We require familiarity with linear algebra, calculus, and statistics.

We will explain what is needed as we go along, but you may need to brush up on some of your old notes, or look it up online.

Now, onto the fun stuff…

What is machine learning?

(patterns)

- *"the **automatic discovery of regularities in data** through the use of computer algorithms and with the use of these regularities to take actions such as classifying the data into different categories"*

  [Bishop, 2007]

- *"Machine learning is the study of computer algorithms that **improve automatically through experience**"*
- *"A computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E."*

  [Mitchell, 1997]

What is machine learning?

- Automating the process that has driven science for centuries
- Collect data, hypothesise relationships (models), experiment to test hypotheses
- E.g.
  - Kepler's laws of planetary motion
  - Newton's laws of motion
  - …

Example: Hand-Written Digits

- Write a program to automatically classify these

Example: Faces

- Write a program to automatically find faces

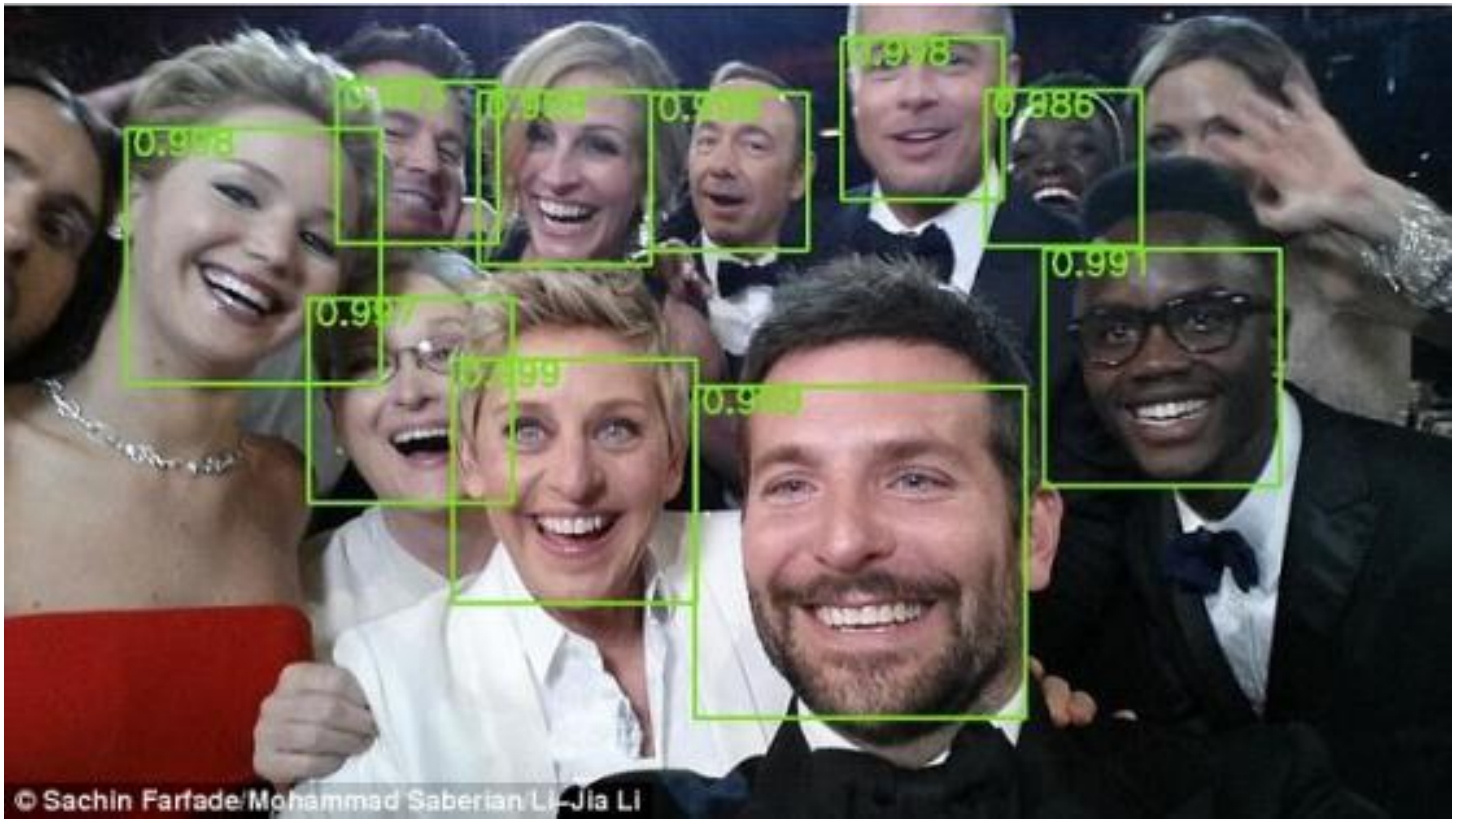

What is machine learning?

- In many of today's problems it is
  - Very hard to write a correct program
    - Heuristics, rules, special cases
    - E.g. based on character strokes, topology, etc
  - But very easy to collect examples
    - Text, emails, images, sales, …
- Idea behind machine learning:
  - From the examples, generate the program (function)
  - This is easier with many examples
    - Define **many** relative to the size of the data, and complexity of the model

Terms

- **Artificial Intelligence (AI)**
- **Machine Learning**
- **Data Science**
- **Big Data**

Terms

- **Artificial Intelligence** (**AI**) is the field which studies how to create computers and computer software that are capable of intelligent behaviour.
- **Machine Learning**
- **Data Science**
- **Big Data**

Terms

- **Artificial Intelligence** (**AI**)
- **Machine Learning** explores the study and construction of algorithms that can learn from and make predictions on data.
- **Data Science**
- **Big Data**

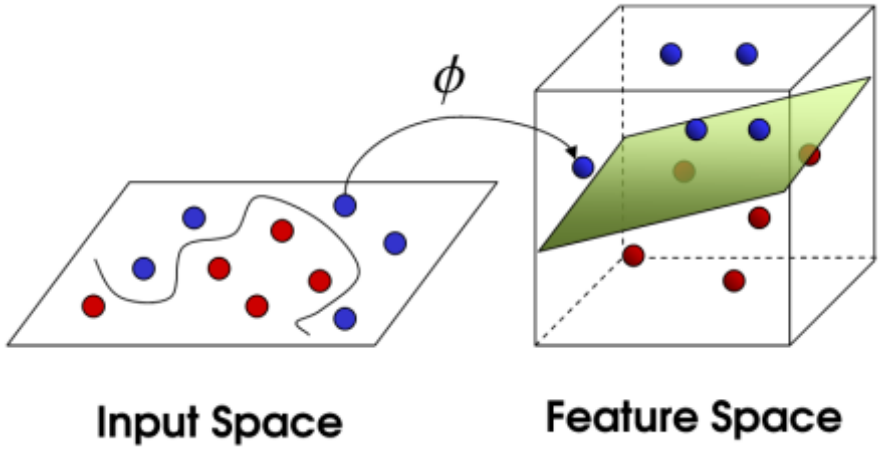

Terms

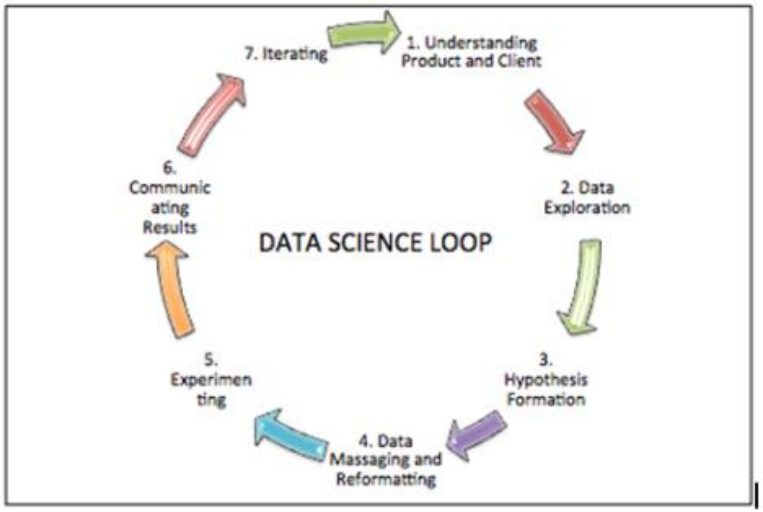

- **Artificial Intelligence** (**AI**)
- **Machine Learning**
- **Data Science** is an interdisciplinary field about processes and systems to extract knowledge or insights from large volumes of data.
- **Big Data**

Terms

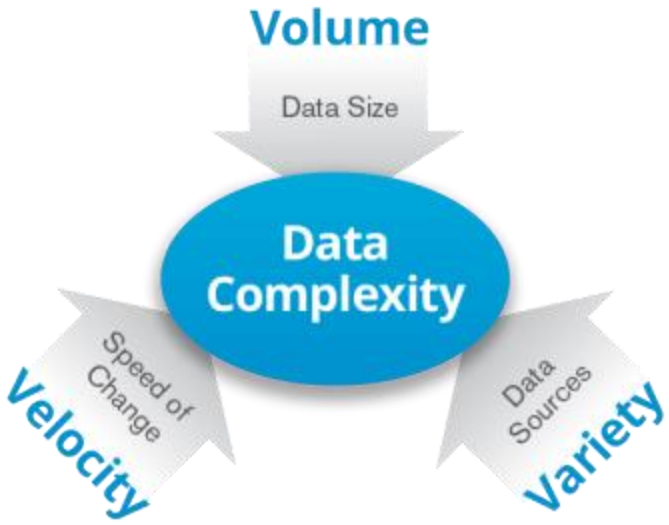

- **Artificial Intelligence** (**AI**)
- **Machine Learning**
- **Data Science**
- **Big Data** is a broad term for data sets so large or complex that traditional data processing applications are inadequate. Challenges include analysis, capture, data curation, search, sharing, storage, transfer, visualization, and information privacy.

ML, Stats and Data Science

Many parts of machine learning have similarities to ideas from statistics

Emphasis is typically different

- Focus on **prediction** in machine learning vs **interpretation** of the model in statistics

ML often refers to tasks associated with artificial intelligence (AI)

- e.g. recognition, diagnosis, planning, robot control, prediction, etc.

Goals can be autonomous machine performance, or enabling humans to learn from data (data mining)

Data science: typically refers to entire pipeline

- From data acquisition and storage to presenting results
- ML is typically the modelling and analysis phase

Typical learning problem

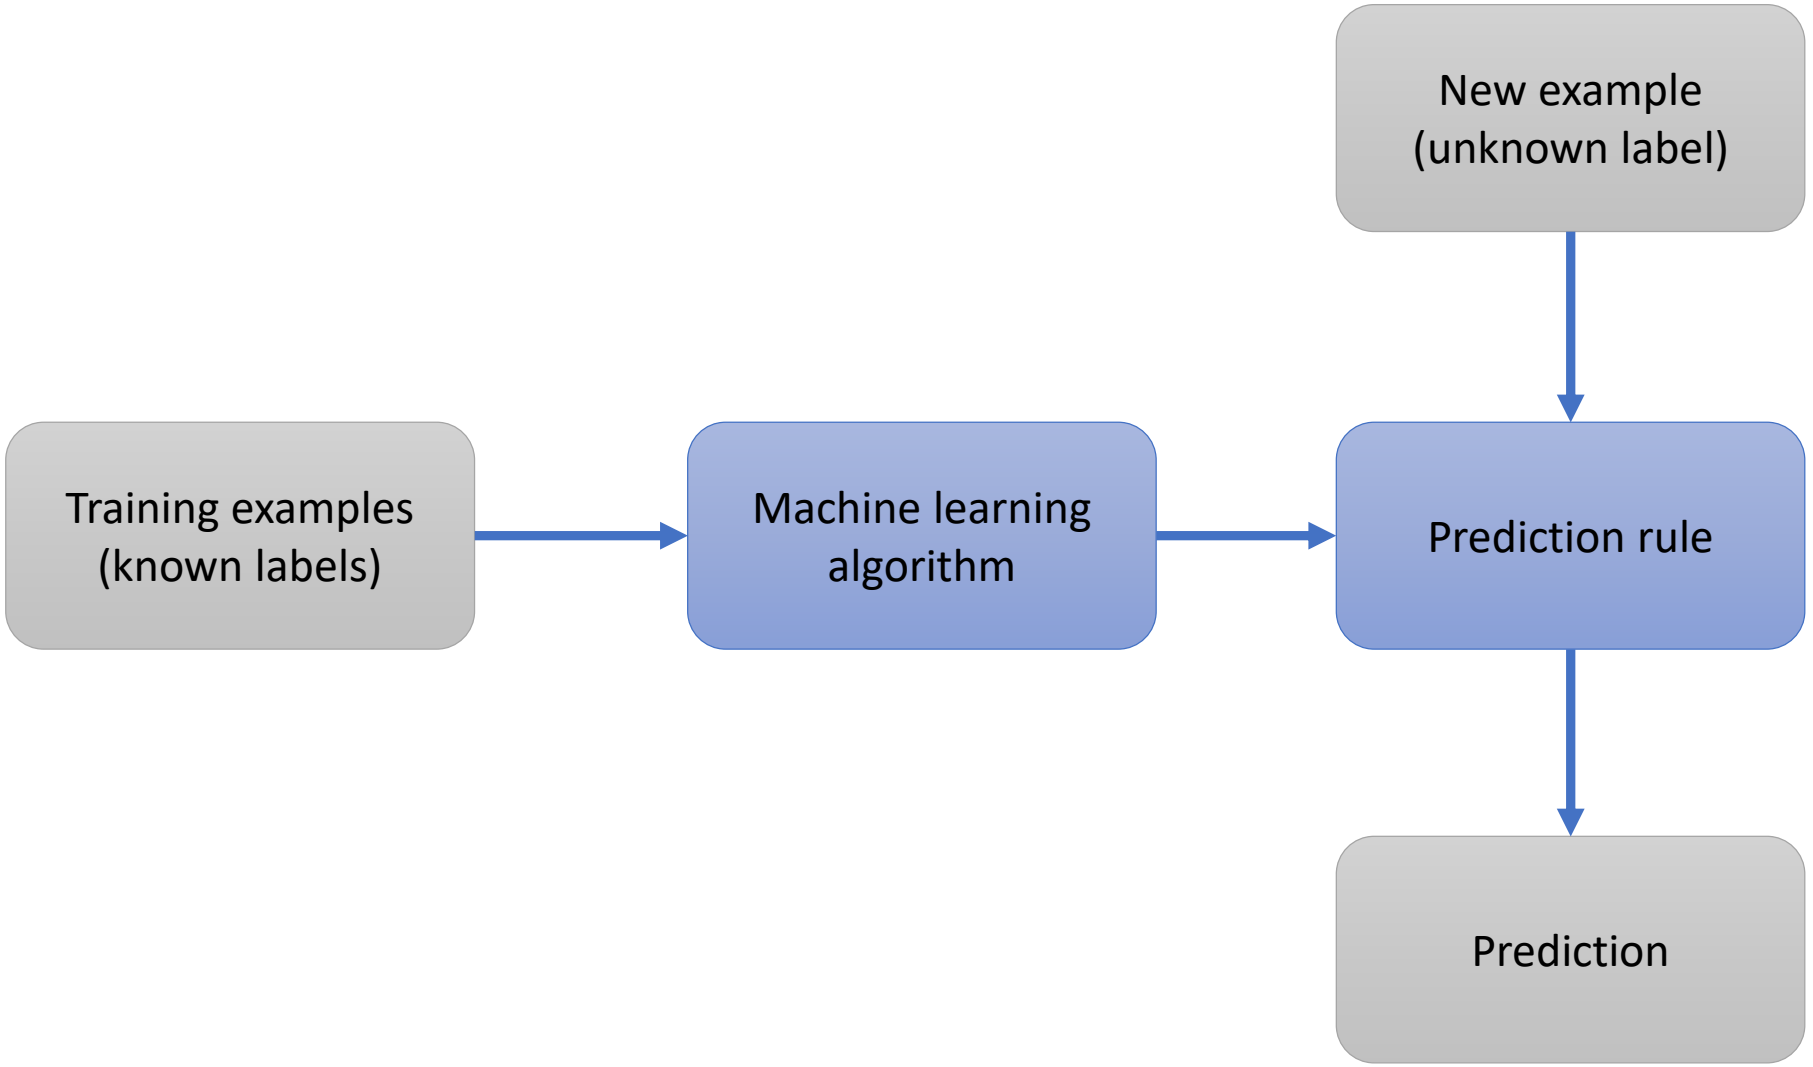

Example problems

- Text recognition, understanding and translation
- Face/object detection
- Spam filtering
- Identifying topics in documents
- Spoken language understanding
- Medical diagnosis
- Customer segmentation / product recommendation
- Fraud detection
- Weather prediction
- Computer game AI

Hand-Written Digits

What is the ML problem here?

Learn a **function** y = f(x):

- Mapping from **data** (x=image) to a **class** (y=number label)
- Think of this as a program
- E.g. f() = “8”

How is this done?

- The function f is some **model** of the digits
  - We provide the general form
- We present the model with a set of N images $\{x_1, x_2, \ldots, x_N\}$ where the true solutions $\{y_1, y_2, \ldots, y_N\}$ are known
- This **training** tweaks the parameters of the model
- During **testing** provide a new x’ to determine y’ = f(x’)

Formally

Given some data (input-output pairs):

- x = training input, y = training output
- D= $\{(x_1, y_1), (x_2, y_2), \ldots, (x_N, y_N)\}$

We need a model:

- *Modelling assumptions go here!*
- $y = f(x; \theta)$
- θ = parameters of model

Now, learning task is to find "best" model parameters θ

- Usually measure some error between y and $f(x; \theta)$
- Treated as an optimisation problem

An example

- Given (features):
  - Animal size
  - Level of domestication
- Predict (label):
  - Cat or dog?

An example

Training

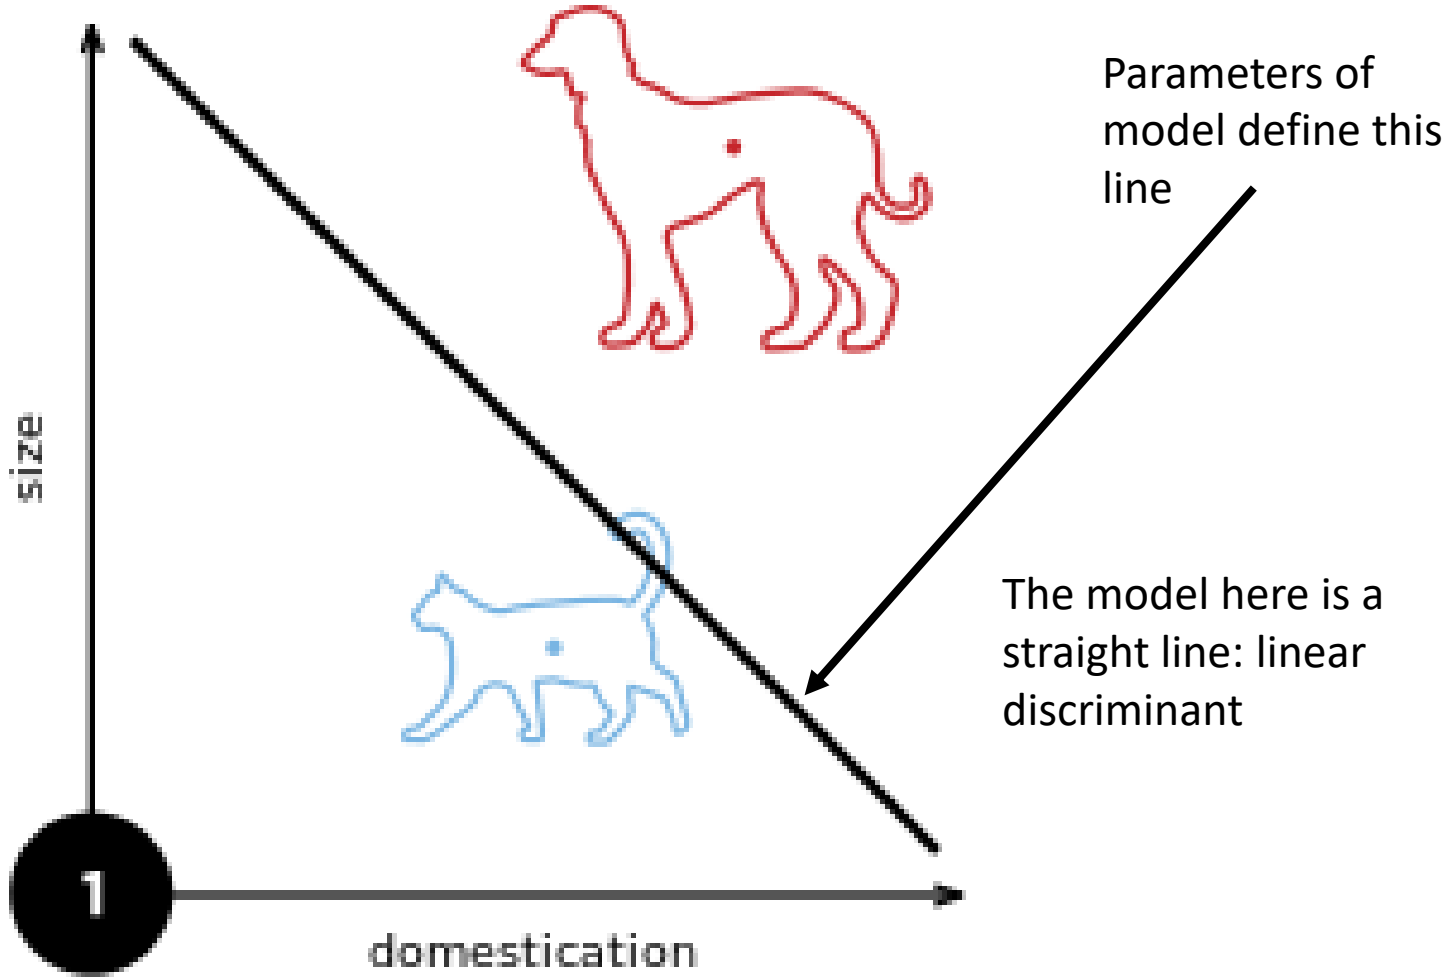

Parameters of model define this line

The model here is a straight line: linear discriminant

An example

Training

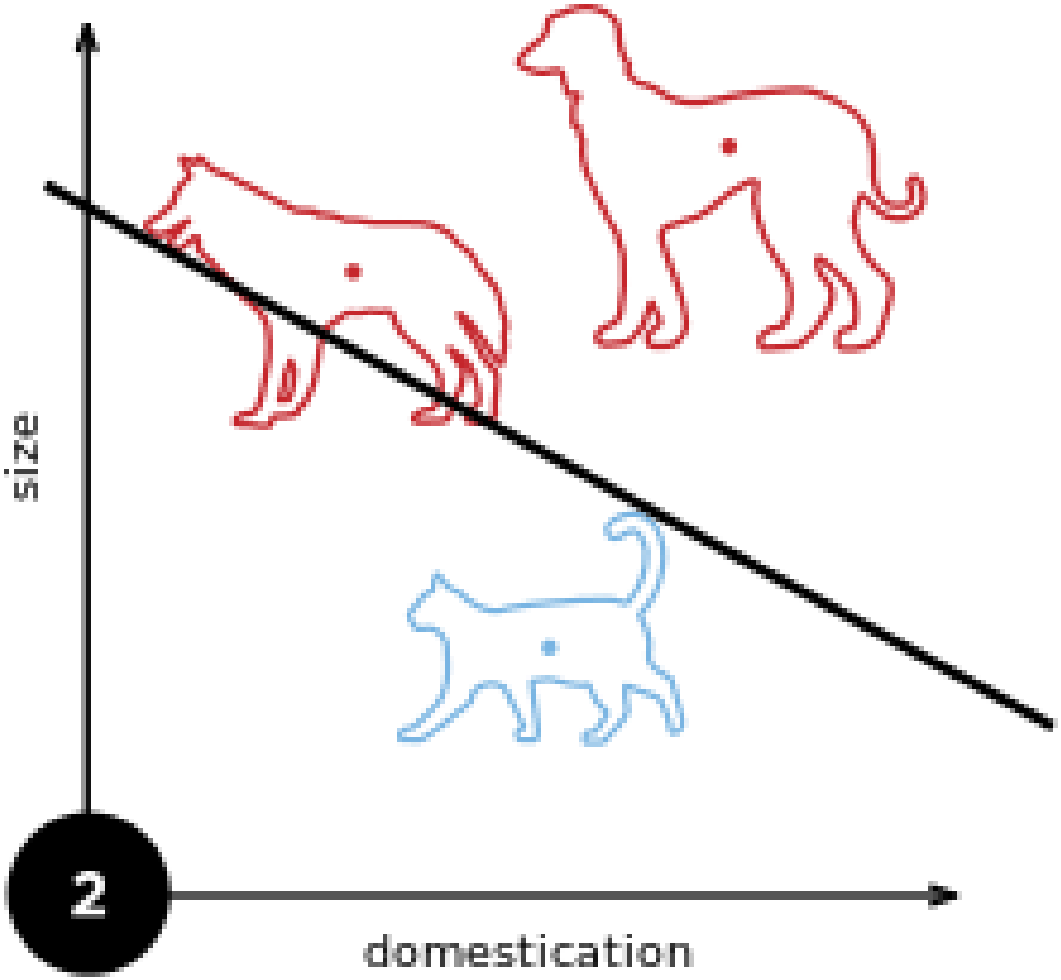

An example

Training

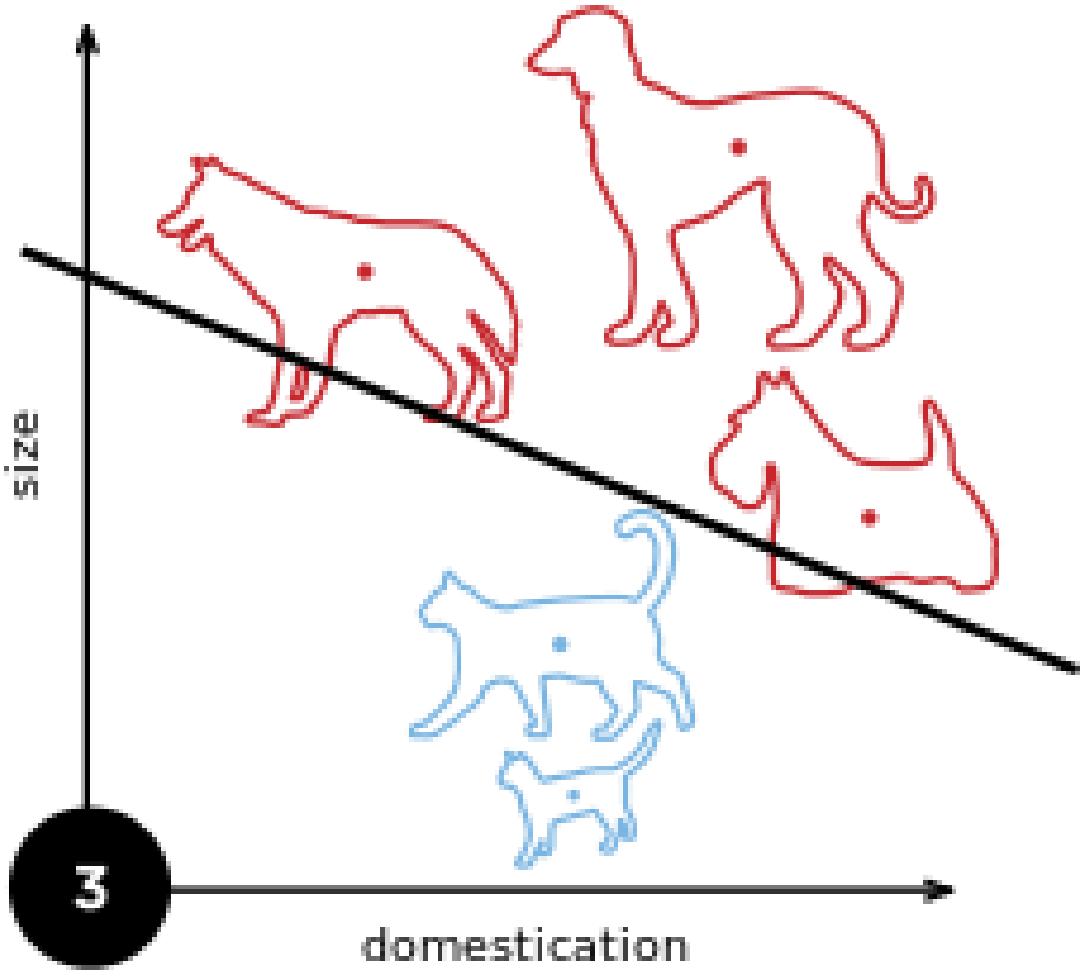

An example

Training

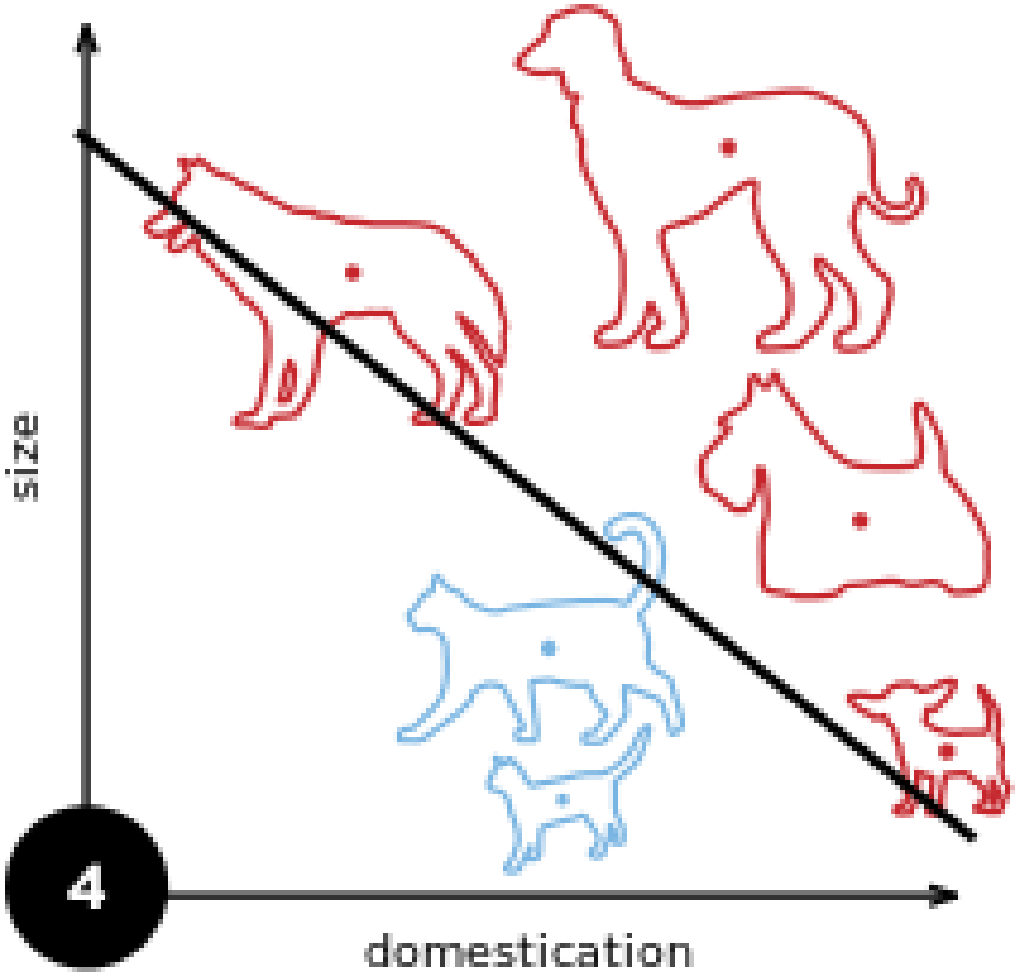

An example

We want the model to **generalise** to any point in this space that we haven't seen yet

Querying

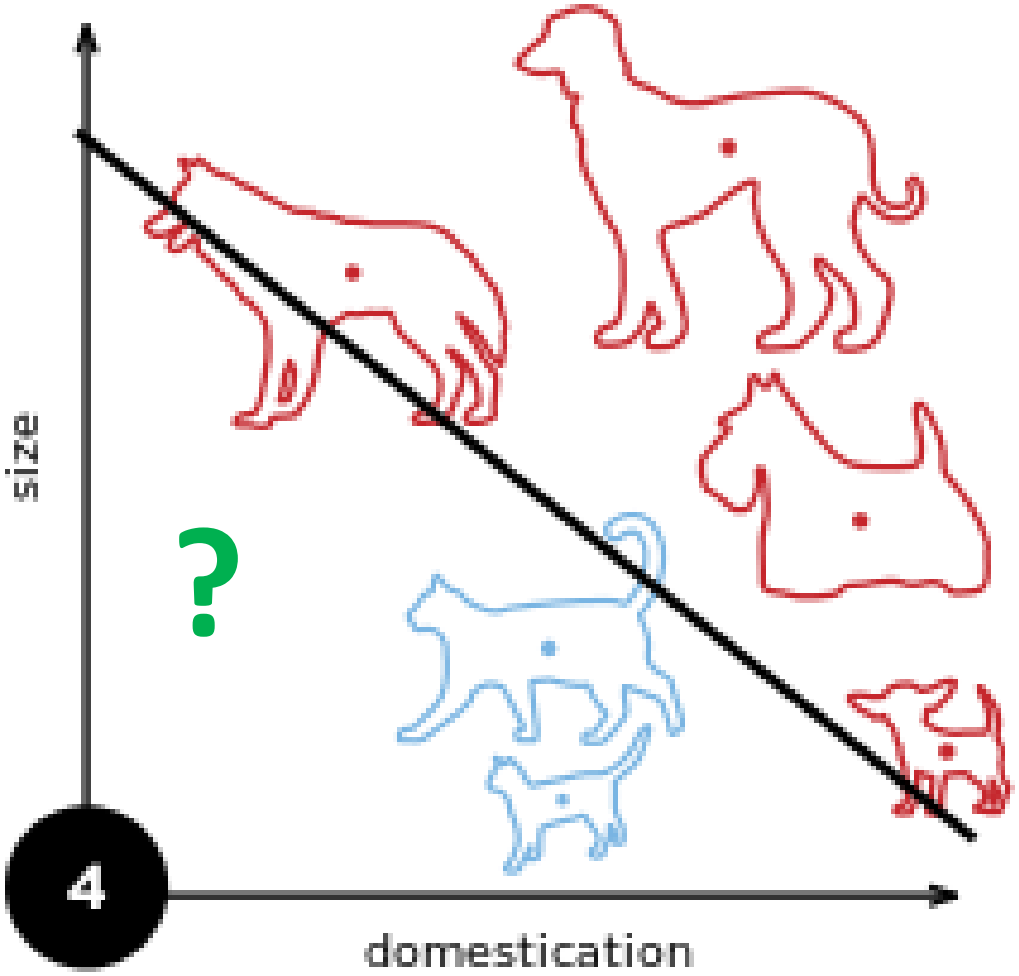

Categories of ML

- Supervised learning
  - Predict output y when given input x
  - Learn from labelled data: $\{(x_i, y_i)\}$
  - Classification: y categorical
  - Regression: y real-valued
- Unsupervised learning
  - Learn from unlabelled data: $\{x_i\}$
  - Clustering
  - Learning some structure in the data
- Semi-supervised learning
  - Only some labels provided
- Reinforcement learning
  - Learn from rewards (typically delayed)
  - Generate own data (experience) through interacting with an environment

Common ML Problems

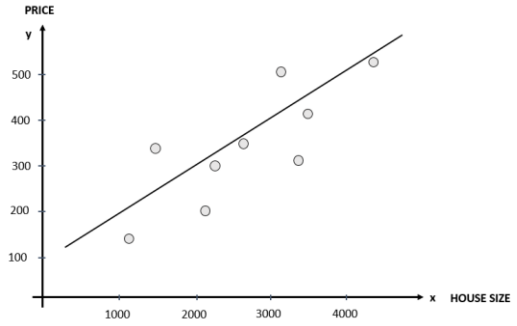

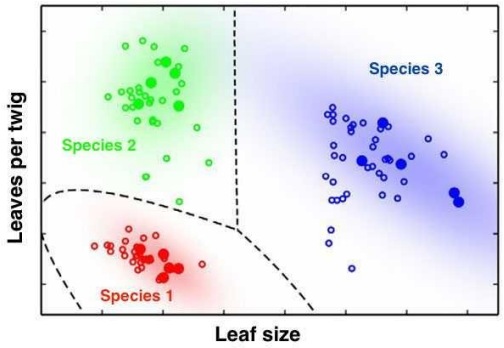

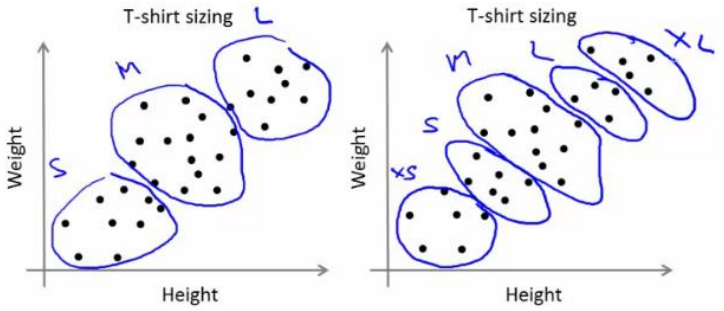

Regression

- Predict a continuous output y, given some input x
- Training: examples of (x,y) pairs
- Supervised learning

Classification

- Predict a discrete class output y, given some input x
- Training: examples of (x,y) pairs
- Supervised learning

Clustering

- Given input data x, return discrete cluster membership
- No training clusters are given, so the returned clusters may not mean anything
- Unsupervised learning

NB: the labels (S, M, L) added to this t-shirt example were the result of someone looking at the clustering results and interpreting them!

Note: inputs are usually multidimensional

Supervised Learning

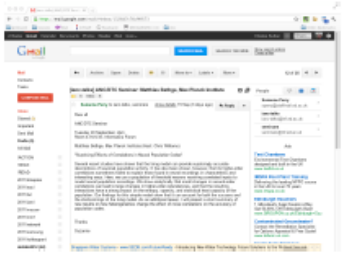

Training data

Feature processing

$\mathbf{x_1} = (1, 0, 0, 3, \ldots.)$ $y_1 = \mathbf{SPAM}$

$\mathbf{x_2} = (-1, 4, 0, 3, \ldots.)$ $y_2 = \mathbf{NOTSPAM}$

Learning algorithm

Classifier

Prediction on new example

$\mathbf{x_{1000}} = (1, 0, 1, 2, \ldots.)$ $y_{1000} = \mathbf{???}$

Unsupervised Learning

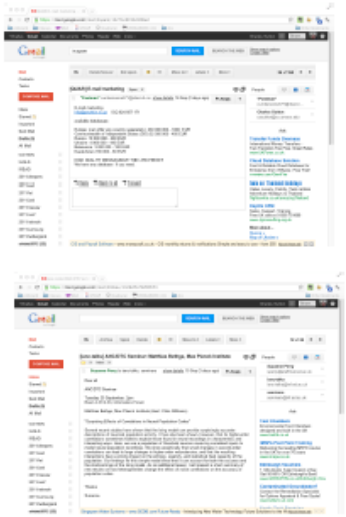

Feature processing

Training data

$\mathbf{x_1} = (1, 0, 0, 3, \ldots.)$

$\mathbf{x_2} = (-1, 4, 0, 3, \ldots.)$

$\ldots.$

$\mathbf{x_{1000}} = (1, 0, 1, 2, \ldots.)$

Learning algorithm

Cluster labels

$c_1 = 4$

$c_2 = 1$

$\ldots.$

$c_2 = 4$

Reinforcement Learning

Given an unknown dynamic process (environment), and an unknown reward process (goal), take a sequence of actions to maximise reward.

Learn a behaviour by rewards received

Many attempts

- Trial and error

Trade-off:

- Exploration:
  - Try something new
- Exploitation:
  - Try the best thing I know

Generalising

During training, only a small fraction of possible data will be provided

Need the resulting f to **generalise** to (ideally) all cases

- This is what we really care about: we are unlikely to only ever see the training data!

Beware of over-/under- fitting!

- Performs well on training data
- Poor generalisation!

The more data the better!

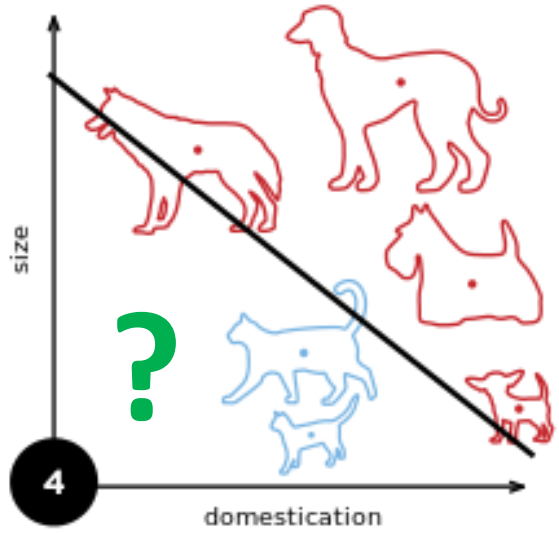

Generalising

Want a model that is expressive, but not too general for the amount of data you have

Regression:

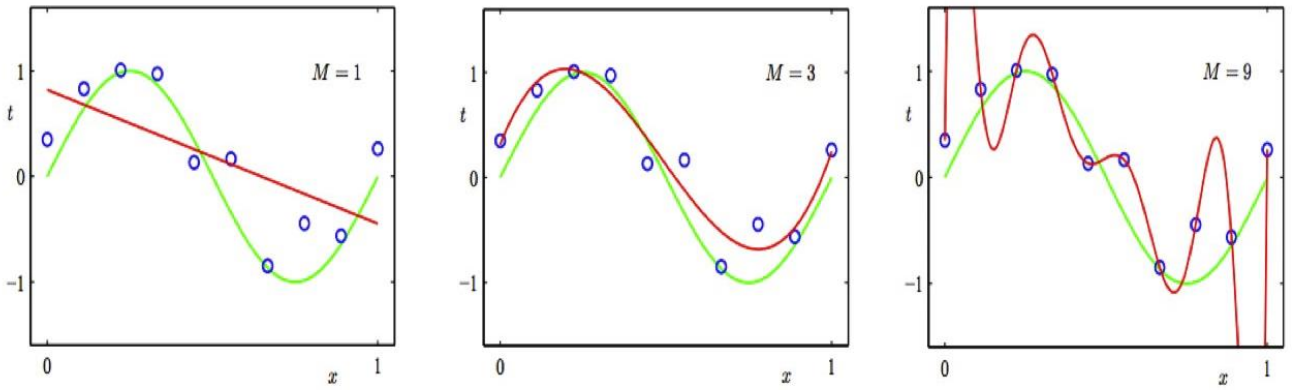

predictor too inflexible: cannot capture pattern

predictor too flexible: fits noise in the data

Note: ground truth (green line) is unknown to algorithm

Classification:

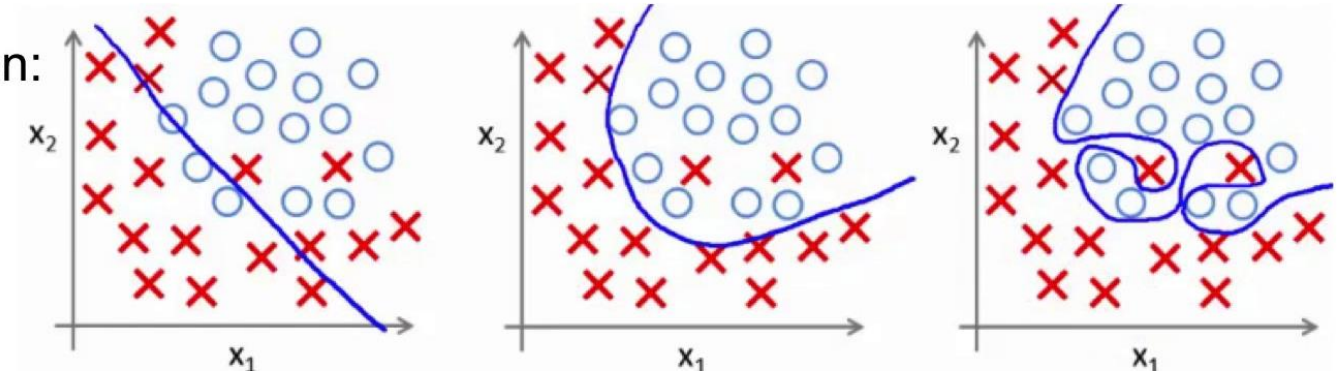

Bias-variance tradeoff

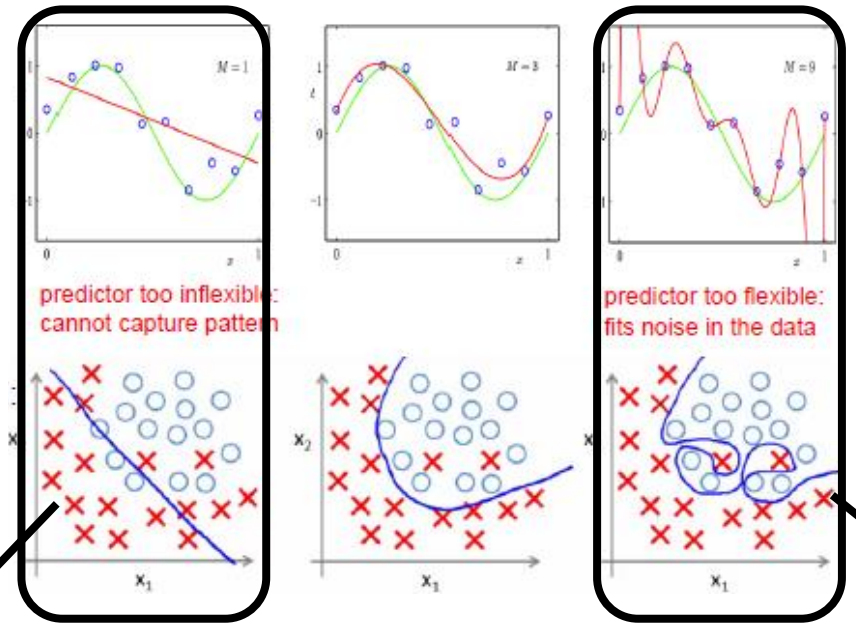

These models are said to have a high **bias**: there is error from erroneous assumptions in the model. This causes **underfitting**.

These models are said to have a high **variance**: there is error from small fluctuations in the data. This causes **overfitting**.

The bias-variance tradeoff involves balancing these two factors: both are different sources of error that affect generalisation.

Representations

How the data is represented is fundamental!

- Determines if problem can be solved with chosen model

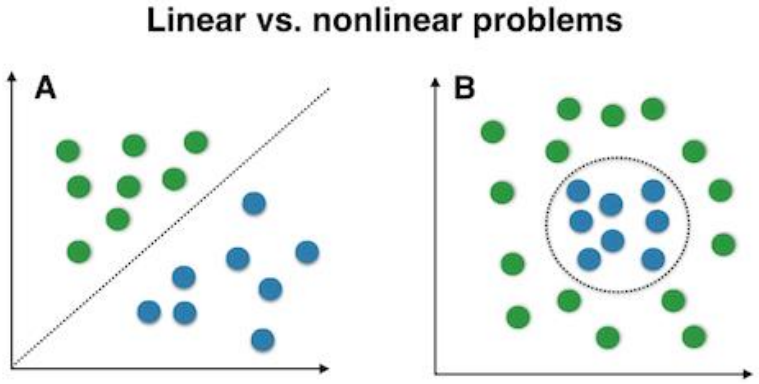

- Or can be solved at all!
  - Identify as elephant vs dog
    1. Features: mass, height
    2. Features: number of legs, number of ears

Representations

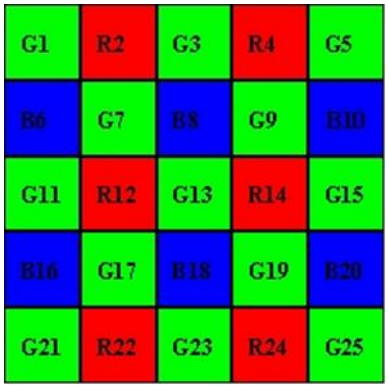

Represent data using features

- Low level: pixels, characters, …
- High level: objects, words, regions, …

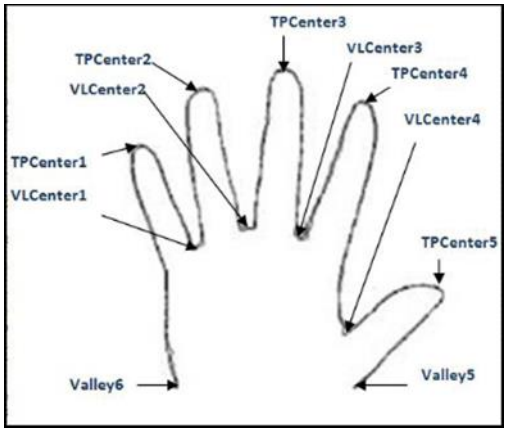

Trade-off between

- Expressive: accurately capture distinctions in data
- Sparse: not need prohibitive amounts of data

One of the hardest and most important parts of ML!

Curse of Dimensionality

- As dimensionality of model or feature space grows, may need **exponentially** more data

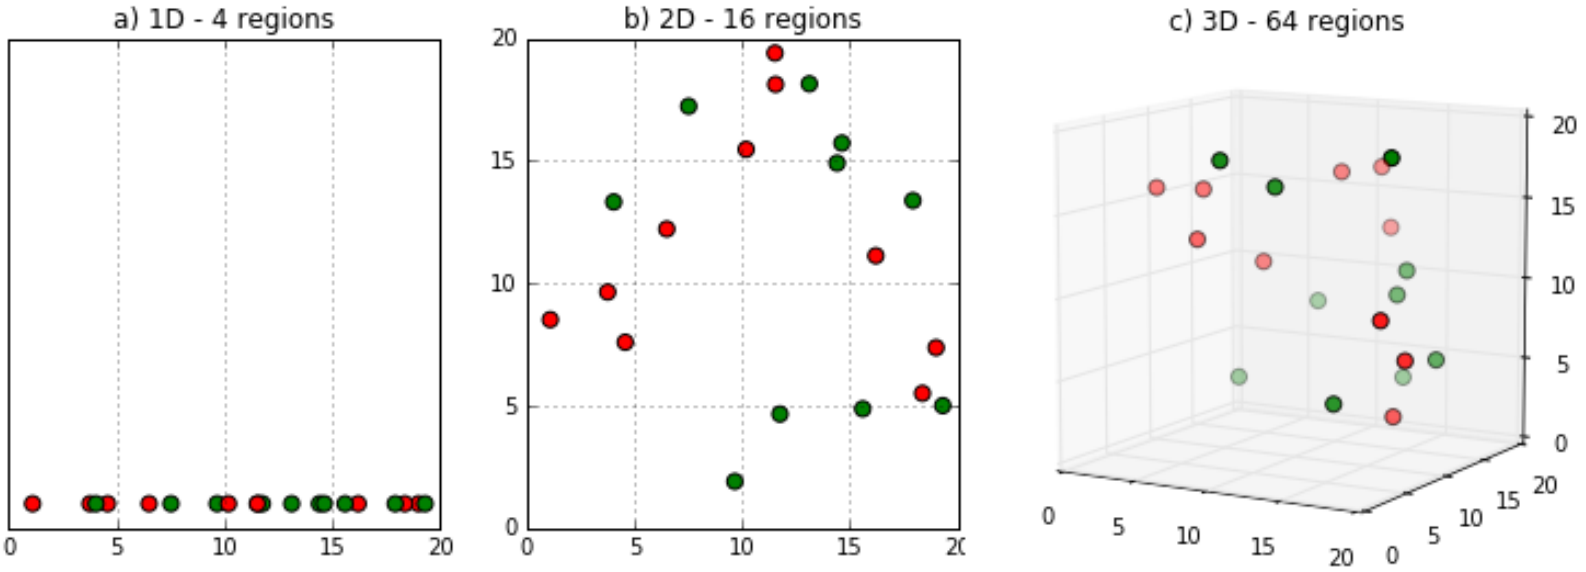

Handling representations

- Feature extraction
  - Manual pre-processing
- Feature selection
  - Autonomously identify important dimensions
- Feature learning
  - Combine simpler features into more complex ones
  - E.g. deep learning (when we talk about neural networks)

Data

For any ML algorithm to work, we need data, and more is always better. In ML, we “let the data do the talking”.

Much work goes into collecting data sets. For large models (many parameters), we may need many millions of examples to learn a good model.

But, how do we know how well

the model will generalise?

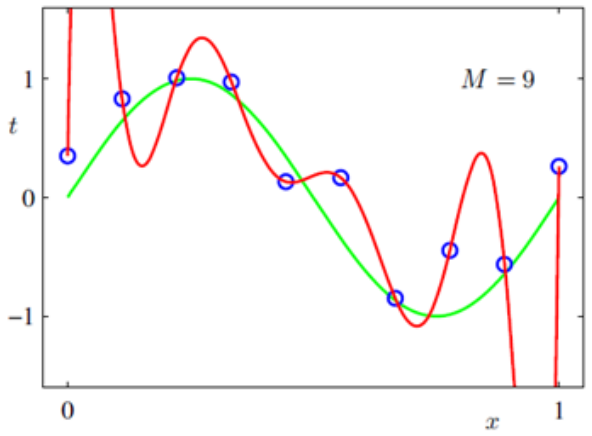

Splitting the Data

Typically divide the full data set into three:

- Training data: learn the model parameters
  - This is the core learning part, and so it needs the most data
  - +/- 60% of the data
- Validation data: learn the model hyperparameters
  - Hyperparameters are values set before training begins, e.g. the degree of the polynomial
  - +/- 20% of the data
- Testing data: report quality of model
  - This is used to report an unbiased evaluation of the final model
  - +/- 20% of the data

Why split the data?

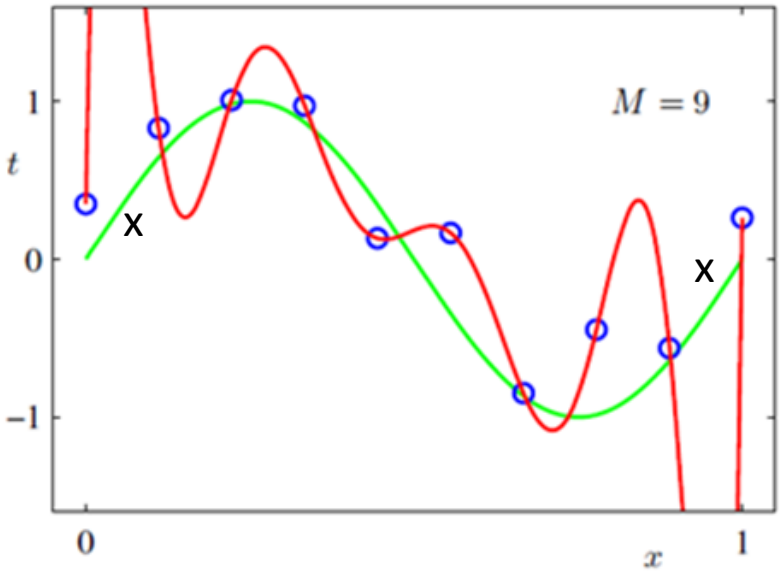

This red model has a perfect fit to the blue training points: so they will not give a reliable estimate of how well the model will generalise.

Instead, we want to **test** it on new data points that **it has never seen during training**. This gives a better idea of its performance.

Similarly, we may be learning the hyperparameter of the degree of the model (M), by training a straight line model (M=1), quadratic model (M=2), … up to M=9 and then seeing which is best. We can train them all on the same training data, but we need to use **different validation data** to choose the best one. Again, we can't just report its performance on that data, as it is already biased. So, we then need a **different testing set** to report final scores.

The test data must not be touched until the very end! It is the "blind/surprise test".

Example: Polynomial Curve Fitting

Simple regression (supervised learning) problem

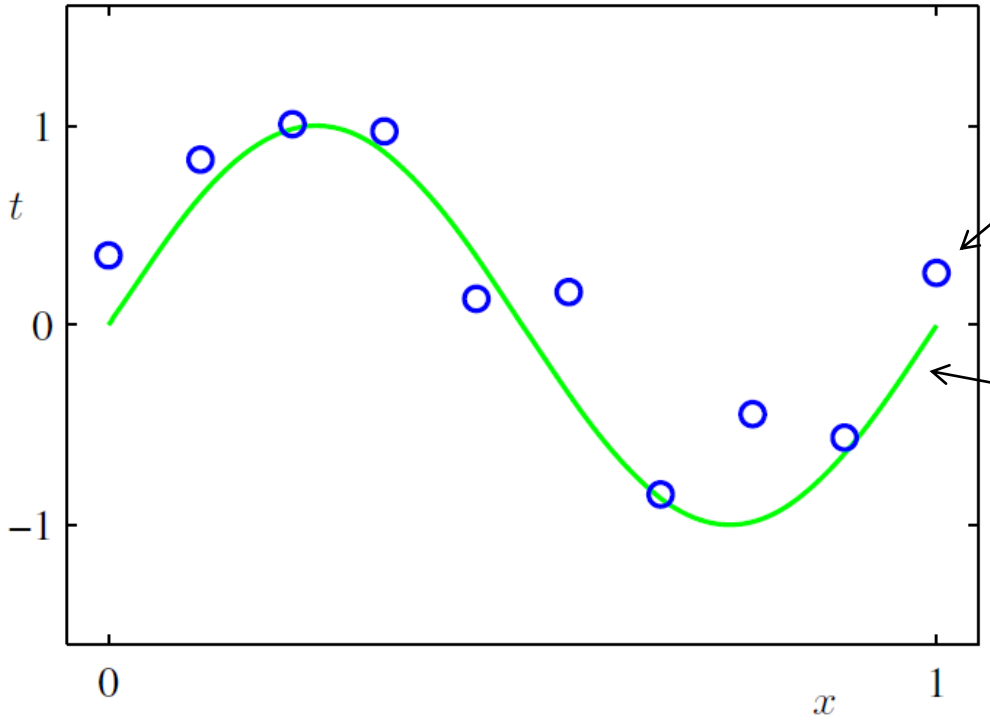

Training data (noisy)

True unknown function: sin(2πx)

Goal: given a new x, predict t (target)

A Polynomial Function

Assume the function is polynomial:

$$y(x, \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + \ldots + w_M x^M = \sum_{j=0}^{M} w_j x^j$$

Note: this is a linear model

- Linear function of coefficients w (the parameters)

Evaluate:

- Use an error function

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^{N} \{y(x_n, \mathbf{w}) - t_n\}^2$$

Learning:

- Find the weight vector w to minimise error E(w)
- Unique solution w*
  - E(w) is quadratic in w
  - E'(w) is linear in w

Predicted value at x
Error between predicted and true value t
Squared so it is symmetrical
Sum over every data point
½ to make the maths simpler after differentiating

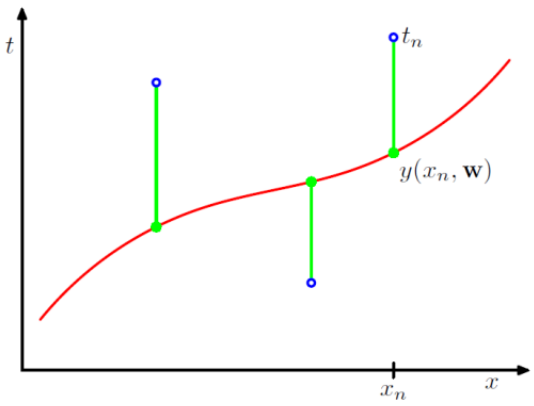

More on this example in the linear regression lecture.

Model Selection

Choosing M (polynomial order)

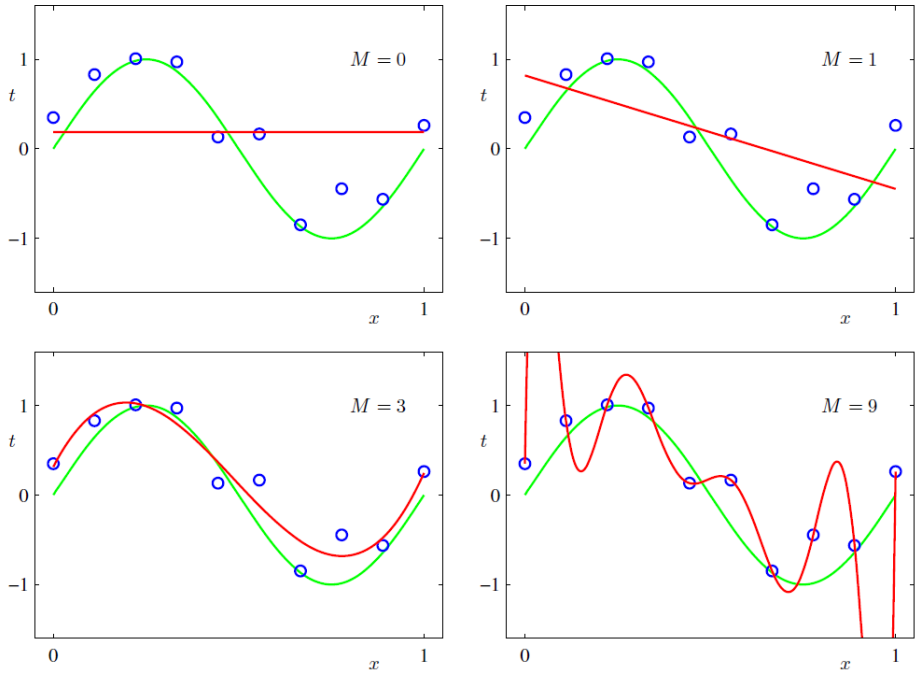

For M = 9, E(w*) = 0! But goal is to **generalise**!

Training vs Testing Error

Define error to compare across N:

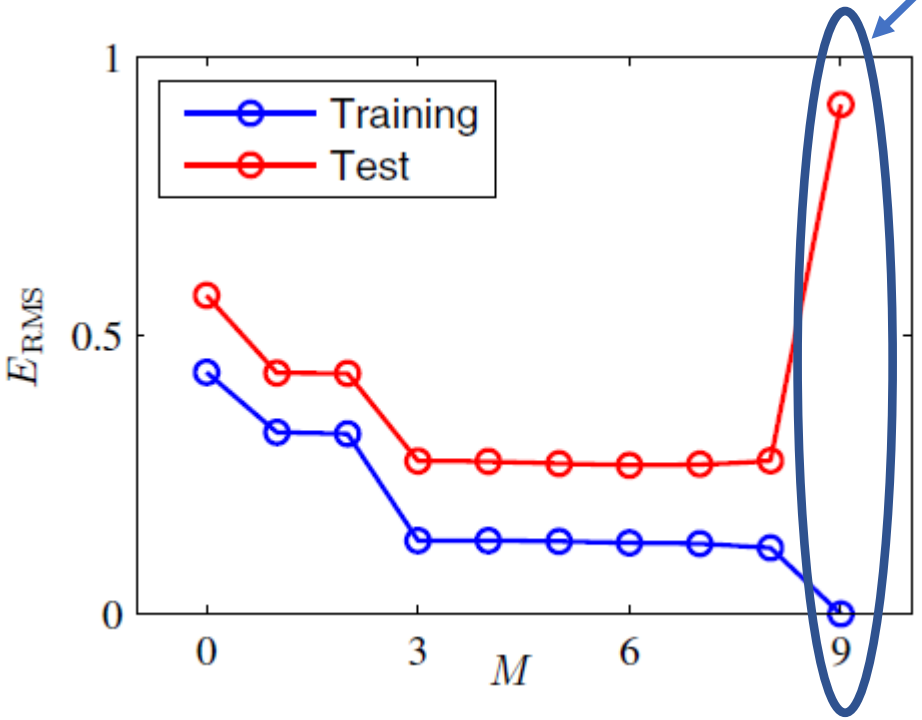

Overfitting: high error on test data, low error on training data

Training error is always better than test error. Why?

Adding More Training Data

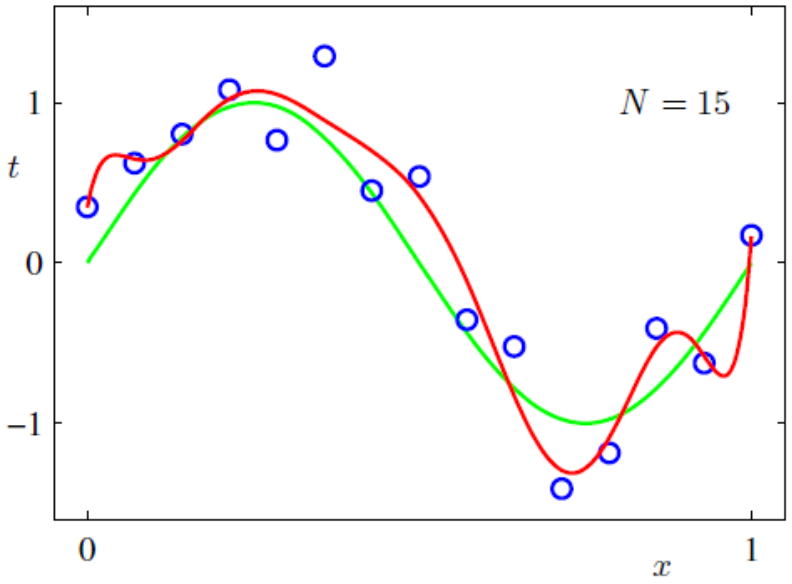

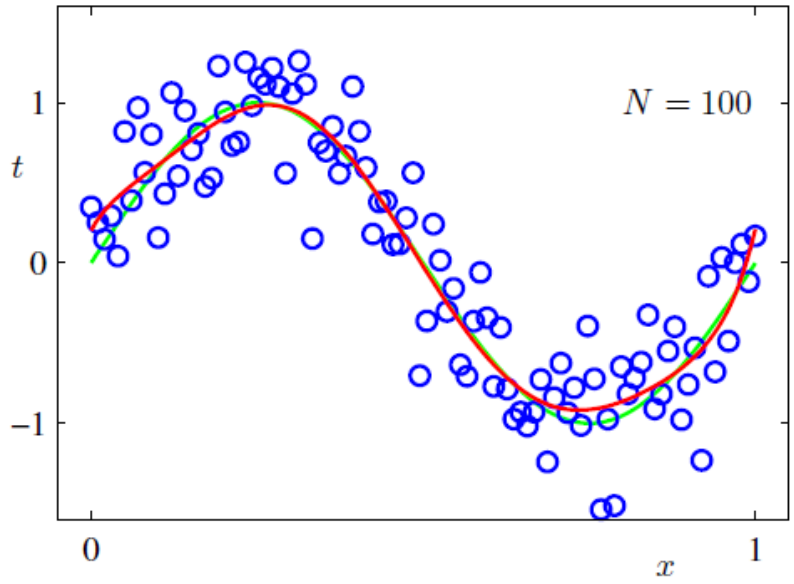

More data

- Less severe over-fitting
- More complex model we can fit

There are other strategies to solve this problem (see regularisation later).

Recap

- What is ML and why do we need it
- Example problems
- Supervised, unsupervised, reinforcement learning
- Generalisation, bias-variance tradeoff
- Representations
- Curse of dimensionality
- Training, validation, testing data
- Curve-fitting example

Make sure you are comfortable with this all by next week!